



**BITS Pilani**  
Pilani | Dubai | Goa | Hyderabad

# Sockets Overview

---

**Srikanth Gunturu**

Guest Faculty  
BITS, WILP

# In this segment

## Sockets Overview

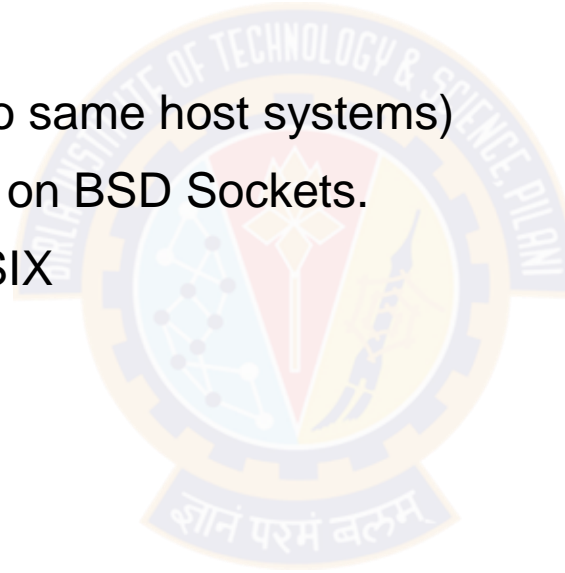
- Introduction
- Data Structures
- Library Calls
- General Operation
- Socket Example



# Socket - Introduction

## What is it ?

- Socket is an internal endpoint for sending/receiving data within a node on network
- Berkeley/POSIX sockets defined API for Inter Process Communication (IPC) within same host (BSD 4.2 – circa 1983)
- Early form of Middleware (limited to same host systems)
- Windows variant (WinSock) based on BSD Sockets.
- Treated similar to files in BSD/POSIX
- Maintained in File Descriptor table
- Supported protocols
  - TCP/IP – IPv4, IPv6
  - UDP



# Socket – Data Structures

## Socket Address


- Defined in *sys/socket.h*
- Address length is 8 bytes by default
- Family – Denotes the family of protocol (TCP, UDP)
- Address contains – host and port
- Socket Address family TCP/IP
  - IPv4 - `sockaddr_in`

```
sa_family_t    sin_family;  
in_port_t      sin_port;  
struct in_addr sin_addr;
```

- IPv6 - `sockaddr_in6`

```
sa_family_t    sin6_family;  
in_port_t      sin6_port;  
uint32_t       sin6_flowinfo;  
struct in6_addr sin6_addr;  
uint32_t       sin6_scope_id;
```

```
struct sockaddr  
{  
    unsigned char  sa_len;    // length of address  
    sa_family_t    sa_family; // the address family  
    char sa_data[14];        // the address  
};
```



# Socket – Library Calls

## Socket APIs

**socket** —creates a descriptor for use in network communications

**connect** —connect to a remote peer (client)

**write** —send outgoing data across a connection

**read** —acquire incoming data from a connection

**close** —terminate communication and deallocate a descriptor

**bind** —bind a local IP address and protocol port to a socket

**listen** —set the socket listening on the given address and port for connections from the client and set the number of incoming connections from a client (backlog) that will be allowed in the listen queue at any one time

**accept** —accept the next incoming connection (server)

**recv** —receive the next incoming datagram

**recvmsg** —receive the next incoming datagram (variation of recv)

**recvfrom** —receive the next incoming datagram and record its source endpoint address

**send** —send an outgoing datagram

**sendmsg** —send an outgoing datagram (variation of send)

**sendto** —send an outgoing datagram, usually to a prerecorded endpoint address

**shutdown** —terminate a TCP connection in one or both directions

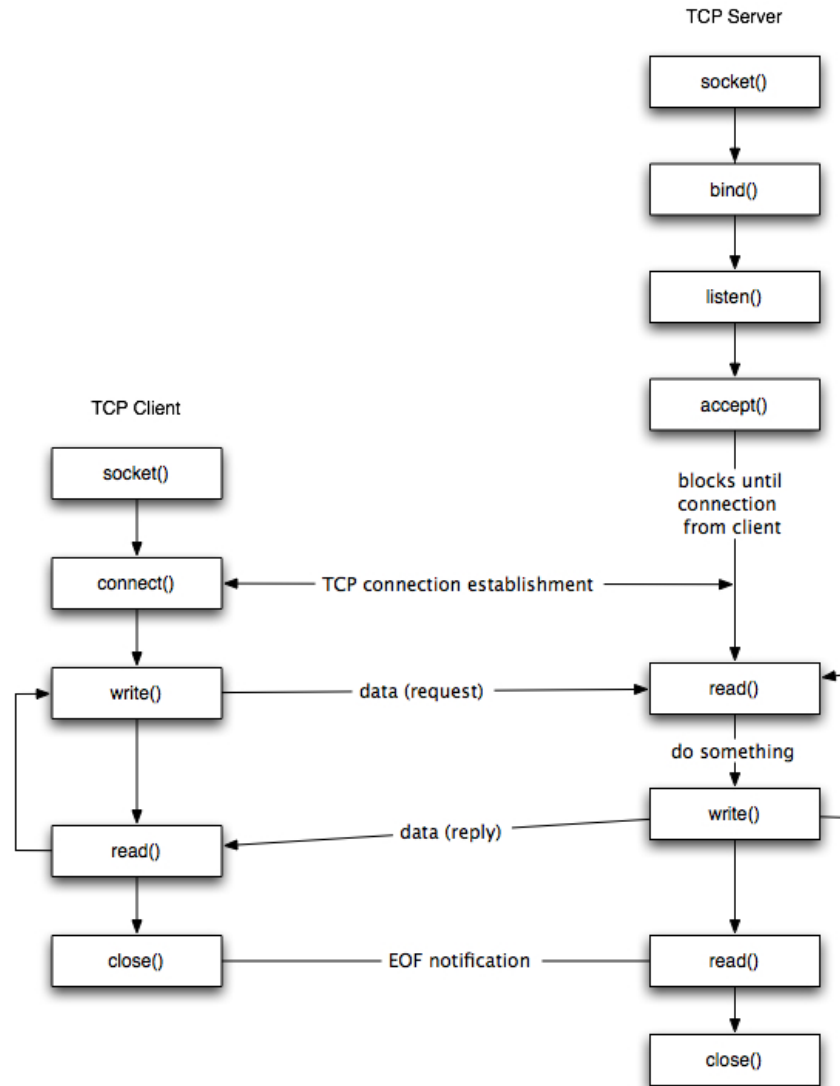
**getpeername** —after a connection arrives, obtain the remote machine's endpoint address from a socket

**getsockopt** —obtain the current options for a socket

**setsockopt** —change the options for a socket

# Socket – General Operation

## Lifecycle





# Socket – Example

## Server code

```
int main()
{
    int server_socket_fd, client_conn_fd;
    int client_addr_size;
    struct sockaddr_in server_addr, client_addr;
    int port_number;
    char message_from_client[256];
    char message_from_server_to_client[256];
    int client_message_length, server_message_length;
    // Set the port number
    port_number = 1132;
    // Create the socket for the server to listen on
    server_socket_fd = socket( AF_INET,
    SOCK_STREAM,
    0
    );
    if (server_socket_fd < 0)
        PrintError("ERROR opening socket");
    // clear out the server address to make sure no problems
    with binding
    // if the address is clear then it won't think the address
    has
    // already been used by another socket
    bzero( (char *) &server_addr, sizeof(server_addr));
```

```
// set the sockaddr in struct appropriately
// note that this assumes IPv4
server_addr.sin_family = AF_INET;
server_addr.sin_addr.s_addr = INADDR_ANY;
server_addr.sin_port = htons(port_number);
// Bind the socket descriptor to the address
if (bind( server_socket_fd,
    (struct sockaddr *) &server_addr,
    sizeof(server_addr))
    < 0)
    PrintError("binding to socket failed");
// Start the server listening on the socket
// limit the number of connections in the listen queue to 3
// (this is the backlog)
listen(server_socket_fd, 3);
while (1)
{
    // clear out the client address to make sure no problems with
    accepting
    // on this address
    // if the address is clear then it won't think the address has
    // already been used by another socket
    bzero( (char *) &client_addr, sizeof(client_addr));
```

# Socket – Example

## Server code

*accept*

```
// accept a connection from a client
// on the open socket
client_addr_size = sizeof(client_addr);
client_conn_fd = accept( server_socket_fd,
(struct sockaddr *) &client_addr,
(socklen_t *) &client_addr_size
);
if (client_conn_fd < 0)
PrintError("the accept failed");
// Clear out the character array to store the client message
// to make sure there's no garbage in it
bzero(message_from_client, 256);
// read the message from the client
client_message_length = read( client_conn_fd,
message_from_client,
255
);
if (client_message_length < 0)
PrintError(" unable to read from socket" );
cout << " message is" << message_from_client << endl;
```

*read*

```
// Now write a message back to the client
// Doing a little mixed mode C++ strings and char buffers
// just to show you how
string mystring;
mystring = " Server received from client, then echoed back to client:" ;
mystring += message_from_client;
bzero(message_from_server_to_client, 256);
mystring.copy(message_from_server_to_client, mystring.length() );
server_message_length = write( client_conn_fd,
(char *) message_from_server_to_client,
strlen(message_from_server_to_client)
);
if (server_message_length < 0)
PrintError(" unable to write to socket" );
sleep(5); // let the client close first, avoids socket address reuse issues
close(client_conn_fd);
} // end while (1)
close(server_socket_fd); // this isn't reached because we used while (1)
// but with a different while loop test condition
// this would be important
return 0;
}
```

*close*



# Socket – Example

## Client code

```
int main()
{
    int socket_fd;
    struct sockaddr_in serv_addr;
    int port_number;
    char * IP_address;
    char client_message[256];
    int message_result;
    // Set the IP address (IPv4)
    IP_address = new char [sizeof(" 127.0.0.1")];
    strcpy(IP_address, " 127.0.0.1"); // could instead have
    copied "localhost"
    // Set the port number
    port_number = 1132;
    // Create the socket
    socket_fd = socket( AF_INET,
    SOCK_STREAM,
    0
    );
    if (socket_fd < 0)
    PrintError(" ERROR opening socket" );
    // clear out the server address to make sure no problems
    with binding
    bzero((char *) &serv_addr, sizeof(serv_addr));
```

IPv4

```
// set the sockaddr_in struct appropriately
serv_addr.sin_family = AF_INET;
serv_addr.sin_port = htons(port_number);
// Use the inet_pton function to convert the IP address to
// binary
if (inet_pton( AF_INET,
IP_address,
&serv_addr.sin_addr
)
< 0
)
PrintError(" Unable to convert IP address to binary to put in
serv_addr" );
// Connect to the server
if (connect( socket_fd,
(struct sockaddr *)
&serv_addr,
sizeof(serv_addr)
)
<0
)
PrintError(" unable to connect to server" );
cout << " Enter message to send to server: " ;
bzero(client_message,256);
```

# Socket – Example

## Client code

```
string mystring;  
getline(cin, mystring); // read the line from standard input  
cout << endl;  
strcpy(client_message, mystring.c_str());  
// Write the message to the socket to send to the server  
message_result = write( socket_fd,  
client_message,  
strlen(client_message)  
);  
if (message_result < 0)  
PrintError("unable to write to socket");  
// Read the return message from the server  
bzero(client_message, 256);  
message_result = read( socket_fd,  
client_message,  
255  
);  
if (message_result < 0)  
PrintError("unable to read from socket");  
cout << client_message << endl;  
// close(socket_fd); // commented out because only close it if  
you  
// don't want to do more calling the server from  
// a run of the client  
delete [] IP_address; // return the memory to the heap  
return 0;  
}
```

Write to  
Socket

read

close



# Thank You!

In our next session:  
Early Middleware Technologies